



## Data wrangling in Python - Data wrangling with Pandas - 2

*One should look for what is and not what he thinks should be. (Albert Einstein)*

# Module completion checklist

Objective	Complete
Define the use and properties of DataFrames, and apply basic operations	
Define the use and properties of Index in DataFrames, and apply basic operations	

# DataFrames

- Now that we have reviewed `Series`, let's look at what a `DataFrame` is
- A **DataFrame** is the single most important object in Pandas
  - It is a collection of `Series` of equal lengths
  - Just like `Series`, `DataFrames` come with many useful methods
- Review complete documentation of the `DataFrame` function [here](#)
- We're now going to build a sample `DataFrame` using two different `Series` to show the average number of days people were out of office in a given month

# Series to DataFrame

- To create a DataFrame, we are going to need to create two Series first
  - The first Series will consist of a range of dates at monthly intervals
  - The second Series will consist of the average number of days people were out of office, in the form of a list of numerical values

```
# Series 1 - times:
times = pd.date_range(start = '20170101', end = '20170630', freq = 'M')

# Series 2 - days out of the office:
days = pd.Series([2, 2, 6, 6, 2, 3])
```

- Now that we have our Series, we can create a DataFrame object with the `pd.DataFrame` function, specifying the two Series we want to include

# Generate DataFrame from Series

- Create a DataFrame using dictionary-like syntax:
  - Dictionary keys become column names of the DataFrame, and
  - Dictionary values become column values
- Inspect the DataFrame by looking at the first few rows, using `.head()`

```
# Create a DataFrame from the two Series we just created, as a dictionary.
average_ooo = pd.DataFrame({'Timestamp': times, '000': days})

# View the first few rows of the DataFrame, using the Pandas function `.head()`.
print(average_ooo.head())
```

	Timestamp	000
0	2017-01-31	2
1	2017-02-28	2
2	2017-03-31	6
3	2017-04-30	6
4	2017-05-31	2

# Look up DataFrame information

- As with arrays and lists, we can look up the `type` of the created object as well as its shape
- A dataframe is a rectangular object, with rows and columns like a matrix:
  - The first number in parentheses gives us the number of rows
  - The second number is the number of columns

```
# Look up the type of object.  
print(type(average_ooo))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
# Look up its shape.  
print(average_ooo.shape)
```

```
(6, 2)
```

# DataFrame description metrics

- There are many metrics you can pull from a `DataFrame` object
- We will now review some key metrics that will help us understand our data
  - `.columns` returns column names
  - `.info()` gives us some extra info about each column like its data type, and how many null values it has
  - `.describe()` computes summary statistics on any numeric column



# DataFrame description metrics (cont'd)

- Now, let's preview these metrics on the average\_ooo dataset

```
print(average_ooo.columns)
```

```
Index(['Timestamp', '000'], dtype='object')
```

```
print(average_ooo.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Timestamp   6 non-null      datetime64[ns]
1   000          6 non-null      int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 224.0 bytes
None
```

```
print(average_ooo.describe())
```

```
count      000
count      6.000000
mean       3.500000
std        1.974842
min        2.000000
25%        2.000000
50%        2.500000
75%        5.250000
max        6.000000
```



# Extracting a single column

- To extract a column, just put its name in quotation marks into square brackets like this:

```
data_frame['column_name']
```

```
print(average_ooo['Timestamp'])
```

```
0    2017-01-31
1    2017-02-28
2    2017-03-31
3    2017-04-30
4    2017-05-31
5    2017-06-30
Name: Timestamp, dtype: datetime64[ns]
```

- The resulting object is a `Series` type
- If you would like to get a `DataFrame` object with a single column, then pass the `list` with a single column name into the square brackets like this:  

```
data_frame[['column_name']]
```

# Extracting multiple columns

- To extract multiple columns, just pass a list of columns

```
print(average_ooo[['Timestamp', 'OOO']])
```

	Timestamp	OOO
0	2017-01-31	2
1	2017-02-28	2
2	2017-03-31	6
3	2017-04-30	6
4	2017-05-31	2
5	2017-06-30	3

# Extracting a single row

- To extract a particular row from a dataframe, we can use a syntax similar to what we used for `ndarray`, but with one small change: **we must use the `iloc` method!**

```
june_ooo = average_ooo.iloc[5, :]  
print(june_ooo)
```

```
Timestamp    2017-06-30 00:00:00  
ooo          3  
Name: 5, dtype: object
```

```
june_ooo = average_ooo.iloc[5] #<- equivalent without the colon  
print(june_ooo)
```

```
Timestamp    2017-06-30 00:00:00  
ooo          3  
Name: 5, dtype: object
```

- `iloc` can be used when the index label of a dataframe is numeric (*integer in `i` loc*) or if you aren't sure of the index label
- Further, we're going to talk more about the importance of the index in Pandas

# Module completion checklist

Objective	Complete
Define the use and properties of DataFrames, and apply basic operations	✓
Define the use and properties of Index in dataframes, and apply basic operations	

# Working with DataFrame indices

- Dataframes in Pandas have a property called the `index`
- The `index` serves many purposes and is an important concept to understand within Pandas
- Some main purposes are:
  - **identifying** data using known indicators, important for analysis, visualization, and interactive console display
  - **enabling** automatic and explicit data alignment
  - **facilitating** intuitive getting and setting of subsets of the dataset



# Index for our dataset

- The **index** of the `average_ooo` dataframe is an **unlabeled column** consisting of the numbers 0 to 5
- By default, the `index` is simply the row number (starting with 0), but it can sometimes make sense to use something more descriptive for the index
- We are going to use `set_index` to set our index in `average_ooo`

```
# Let's use the `Timestamp` column as our new index.  
average_ooo = average_ooo.set_index('Timestamp')  
print(average_ooo)
```

Timestamp	ooo
2017-01-31	2
2017-02-28	2
2017-03-31	6
2017-04-30	6
2017-05-31	2
2017-06-30	3

# Looking up by the new index

- Now the rows of our dataframe are indexed by the time stamp and the `Timestamp` column has been removed
- This makes it really easy to look up values corresponding to a particular time stamp, rather than to an arbitrary row number
- To do so, we now use the `.loc()` method

```
print(average_ooo.index)
```

```
DatetimeIndex(['2017-01-31', '2017-02-28', '2017-03-31', '2017-04-30',  
              '2017-05-31', '2017-06-30'],  
              dtype='datetime64[ns]', name='Timestamp', freq=None)
```

```
# Look up a specific row by index.  
print(average_ooo.loc['2017-02-28'])
```

```
ooo      2  
Name: 2017-02-28 00:00:00, dtype: int64
```

# Loc vs. iloc

- Notice that after implementing our new index, we used `loc` rather than `iloc`
- The “i” in `iloc` stands for integer, so we can only pass it numerical values
- `iloc` is helpful if
  - Your index still consists of the default integer values, i.e. the row numbers
  - You need to reference your rows by their numeric order and not by the index value you’ve set (e.g. by using a counter in a loop)
- Since the row we wanted was in position 1, we could also use `iloc` to look it up:

```
print(average_ooo.iloc[1])
```

```
ooo      2  
Name: 2017-02-28 00:00:00, dtype: int64
```



# Reset the index

- We can change the index back to the default integer values by using `.reset_index()`
- Note that the original unlabeled index column is back and that `Timestamp` is back in the `DataFrame` as a labeled column

```
average_ooo = average_ooo.reset_index()
print(average_ooo.index)
```

```
RangeIndex(start=0, stop=6, step=1)
```

```
# You can see that now `Timestamp` is once again a column vs. what it looked like when it was an
index.
print(average_ooo)
```

	Timestamp	ooo
0	2017-01-31	2
1	2017-02-28	2
2	2017-03-31	6
3	2017-04-30	6
4	2017-05-31	2
5	2017-06-30	3

# Knowledge check



Link: <https://forms.gle/rLRZXU8MDaTUwxSc6>

# Module completion checklist

Objective	Complete
Define the use and properties of DataFrames, and apply basic operations	✓
Define the use and properties of Index in DataFrames, and apply basic operations	✓

# Congratulations on completing this module!

You are now ready to try Tasks 5-8 in the Exercise for this topic

